

Assignment 2

Devansh Jain, 190100044

26 Sept 2021

Contents

1	Perceptron	1
1.1	CS 337: Theory	1
1.	1
2.	1
3.	1
4.	2
5.	2
1.2	CS 335: Lab	2
2	LASSO and ISTA	3
2.1	CS 337: Theory	3
1.	3
2.	3
3.	4
2.2	CS 337: Lab	4
1.	4
2.	5
3.	5
3	Bias Variance Trade-off	6
3.1	CS 337: Theory	6
1.	6
2.	6
3.2	CS 337: Lab	6
1.	6
2.	6
3.	7

1 Perceptron

1.1 CS 337: Theory

1.

Both 1-vs-1 and 1-vs-rest approach use binary classification algorithm.

1-vs-rest has K classifiers - one for each class.

1-vs-1 has $K(K-1)/2$ classifiers - one for each pair of classes.

In 1-vs-rest, each classifier predicts score (probability) and the class with highest score is chosen.

In 1-vs-1, each classifier predicts one class and the class which has been predicted the most is chosen.

Advantages and Disadvantages

Single classifier in 1-vs-1 uses subset of data, so each classifier is faster for 1-vs-1.

1-vs-rest trains less number of classifiers and hence is faster overall.

1-vs-1 is less prone to imbalance in dataset due to dominance in specific classes.

2.

$$\begin{aligned}
 score(f, y) &= \sum_i f_i w_{y_i} \\
 &= f^T w_y \\
 score(f, y)_{new} &= f^T w_{y_{new}} \\
 &= f^T w_{y_{old}} + f^T f \\
 &= f^T w_{y_{old}} + ||f||_2^2 \\
 &\geq f^T w_{y_{old}} \\
 &\geq score(f, y)_{old} \\
 score(f, y')_{new} &= f^T w_{y'_{new}} \\
 &= f^T w_{y'_{old}} - f^T f \\
 &= f^T w_{y'_{old}} - ||f||_2^2 \\
 &\leq f^T w_{y'_{old}} \\
 &\leq score(f, y')_{old}
 \end{aligned}$$

Hence, proved.

3.

$$\begin{aligned}
 \mathcal{L}(f, y) &= \max(0, -yf^T w) \\
 \nabla_w \mathcal{L}(f, y) &= \begin{cases} 0 & \text{if } yf^T w \geq 0 \\ -yf & \text{if } yf^T w < 0, \text{ i.e. point is misclassified} \end{cases}
 \end{aligned}$$

Gradient descent update:

$$\begin{aligned}
 w^{(k+1)} &= w^{(k)} - \eta \nabla_w \mathcal{L}(f, y) \\
 w^{(k+1)} &= \begin{cases} w^{(k)} & \text{if } yf^T w \geq 0 \\ w^{(k)} + \eta yf & \text{if } yf^T w < 0, \text{ i.e. point is misclassified} \end{cases}
 \end{aligned}$$

We can see that the gradient descent update rule is similar to perceptron update rule. Using the fact that the perceptron update rule converges for linearly separable dataset, we can conclude that gradient descent algorithm also converges for linearly separable dataset. Hence, proved.

P.S. The loss function takes into account only a single data point. To be precise, we are doing stochastic gradient descent here.

4.

In proof of the convergence theorem, if we use $\eta = 0.5$.

We get $\sqrt{kr} \geq \|w^{(k)}\|_2 \geq u^T w^{(k)} \geq k\eta\gamma$.

Thus, $k \leq \frac{r^2}{\eta^2\gamma^2}$.

Upper bound on number of iterations under the modified algorithm is 4M.

Alt:

If we do double update at every step, i.e. take the same point again for every misclassification. It would be equivalent to original perceptron, so number of iterations would be 2M.

This is possible only if we chose the points in this fashion.

5.

$k \leq \frac{r^2}{\gamma^2}$, where $r \geq \|f\|_2 \forall f \in \mathcal{D}$ and $\exists u, \|u\|_2 = 1, \gamma \leq |u^T f| \forall f \in \mathcal{D}$.

Here, $f = [f_1 \ f_2 \ 1]$.

So, $r = \sqrt{3}$ and if we take $u = \frac{[1 \ 1 \ -1.5]}{\sqrt{4.25}}$, then $\gamma = \frac{1}{\sqrt{17}}$.

Thus, $k \leq 51$.

1.2 CS 335: Lab

Code for the function `perceptron_algo()` updated in notebook.

	precision	recall	f1-score	support
class 0	1.00	1.00	1.00	13
class 1	1.00	1.00	1.00	14
class 2	1.00	1.00	1.00	13
accuracy			1.00	40
macro avg	1.00	1.00	1.00	40
weighted avg	1.00	1.00	1.00	40

```
Out[17]: array([[ 9.98695753, 11.8197221 ],
                [-15.3119201,  9.86645211],
                [ 5.32496257, -21.68617421]])
```

Figure 1: Classification Report

2 LASSO and ISTA

2.1 CS 337: Theory

1.

Given $y_i = x_i \cdot w + \epsilon_i$, where $\epsilon_i \sim^{iid} \mathcal{N}(0, \sigma^2)$ and $w_j \sim^{iid} \text{Lasso}(0, \theta)$.

$$\begin{aligned}\mathbb{P}(\mathcal{D}|w) &\propto \exp\left(-\frac{\sum_{i=1}^n \|y_i - x_i \cdot w\|_2^2}{2\sigma^2}\right) \\ \mathbb{P}(w) &\propto \exp\left(-\frac{\sum_{j=1} |w_j|}{\theta}\right) \\ &\propto \exp\left(-\frac{\|w\|_1}{\theta}\right) \\ \mathbb{P}(w|\mathcal{D}) &\propto \exp\left(-\frac{\sum_{i=1}^n \|y_i - x_i \cdot w\|_2^2}{2\sigma^2} - \frac{\|w\|_1}{\theta}\right)\end{aligned}$$

$$\begin{aligned}w_{\text{MAP}} &= \arg \max_w \mathbb{P}(w|\mathcal{D}) \\ &= \arg \max_w \exp\left(-\frac{\sum_{i=1}^n \|y_i - x_i \cdot w\|_2^2}{2\sigma^2} - \frac{\|w\|_1}{\theta}\right) \\ &= \arg \min_w \left(\frac{\sum_{i=1}^n \|y_i - x_i \cdot w\|_2^2}{2\sigma^2} + \frac{\|w\|_1}{\theta}\right) \quad (-\log(x) \text{ is a non-increasing function}) \\ &= \arg \min_w \sum_{i=1}^n \|y_i - x_i \cdot w\|_2^2 + \lambda \|w\|_1 \quad \left(\lambda := \frac{2\sigma^2}{\theta}\right) \\ &= \arg \min_w L(w)\end{aligned}$$

Hence, proved.

2.

Referencing from the lecture slides and book by Tibshirani:

As shown in Figure 2, the contours of the error and constrained functions meet at axes for l_1 norm unlike l_2 norm.

This results in sparser weight vector as weights which are too closed to zero in OLS and Ridge are made zero here.

We can also see that l_2 norm penalizes larger values of weights much more than l_1 norm, so in l_2 norm, having a small non-zero value is more likely than l_1 norm.

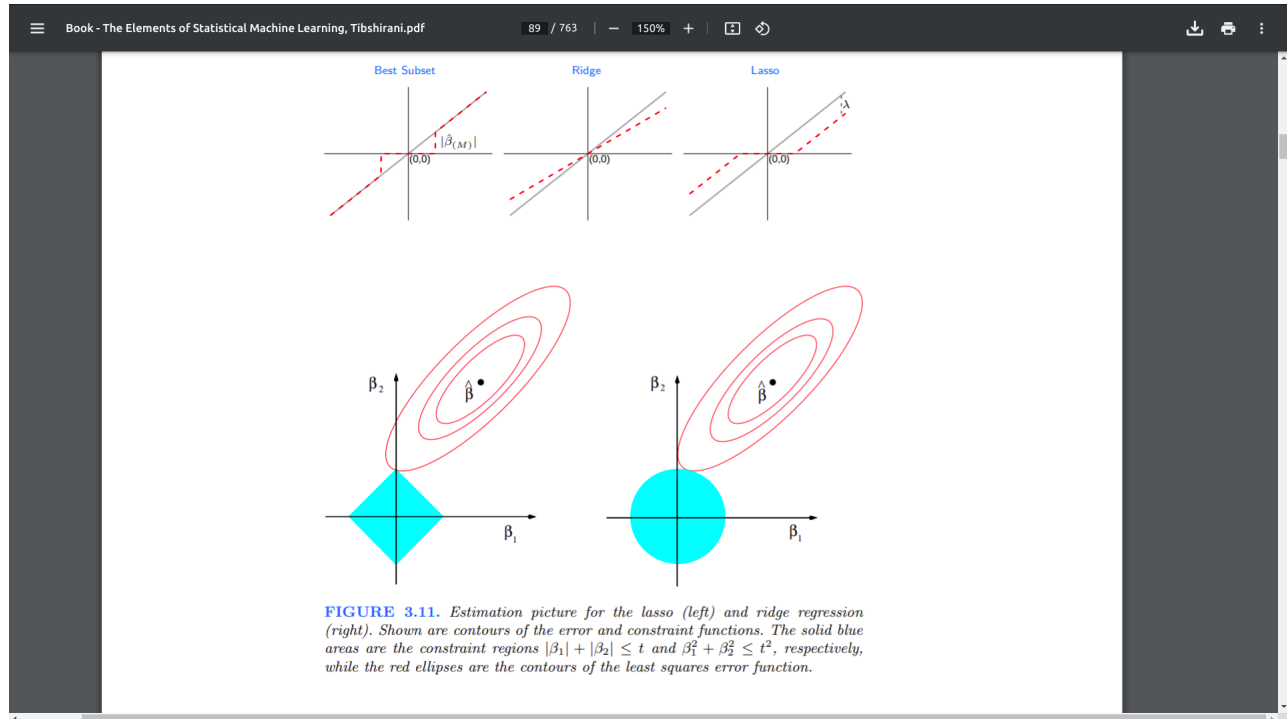


Figure 2: Tibshirani Book: Page 71

3.

Yes, a closed form for Lasso exists when the features are uncorrelated.

We can see this from OLS also, there was a unique solution only when the features were linearly independent.

Same way here, if the features are uncorrelated then there is a unique solution which can be found using subgradient methods (Referred to books and search).

Orthonormal covariates [\[edit\]](#)

Some basic properties of the lasso estimator can now be considered.

Assuming first that the covariates are [orthonormal](#) so that $(x_i | x_j) = \delta_{ij}$, where $(\cdot | \cdot)$ is the [inner product](#) and δ_{ij} is the [Kronecker delta](#), or, equivalently, $X^T X = I$, then using [subgradient methods](#) it can be shown that

$$\hat{\beta}_j = S_{N\lambda}(\hat{\beta}_j^{\text{OLS}}) = \hat{\beta}_j^{\text{OLS}} \max \left(0, 1 - \frac{N\lambda}{|\hat{\beta}_j^{\text{OLS}}|} \right) \quad [2]$$

where $\hat{\beta}^{\text{OLS}} = (X^T X)^{-1} X^T y$

S_α is referred to as the soft thresholding operator, since it translates values towards zero (making them exactly zero if they are small enough) instead of setting smaller values to zero and leaving larger ones untouched as the hard thresholding operator, often denoted H_α , would.

Figure 3: Wikipedia: Lasso (statistics)

2.2 CS 337: Lab

1.

Code for the function `ista()` updated in notebook.




Figure 4: LASSO Regression - ISTA (epochs=10000, lr=1e-2, lamda=0.01)

2.

Code for the functions `mse_multi_var()`, `mse_regularized()`, `split_data()`, `multivar_reg_closedform()` updated in notebook.

3.

If we take `lamda=0.5`, then the second and fourth features (index=1, 3) have weights close to zero for Ridge and exactly zero for LASSO.

This shows that LASSO is able to completely ignore relatively redundant features.

In other words, LASSO can be used for feature selection.

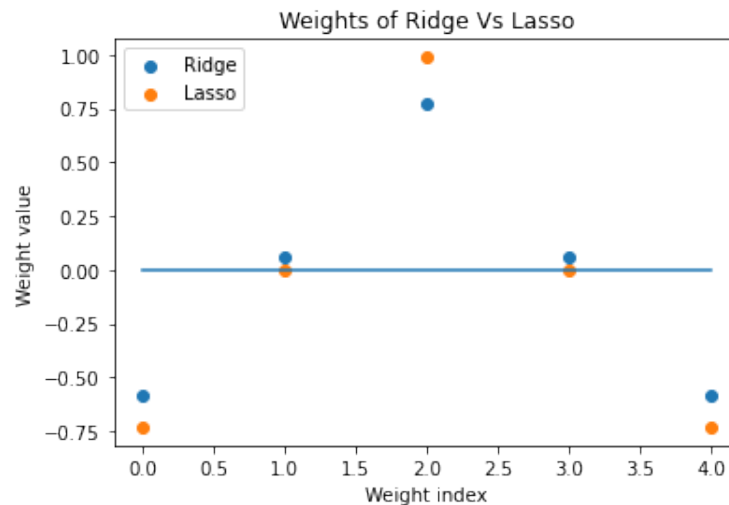


Figure 5: Weights for Ridge and LASSO (`lamda=0.5`)

3 Bias Variance Trade-off

3.1 CS 337: Theory

1.

As discussed in class, Model complexity increases \implies Bias decreases and Variance increases and, Model complexity decreases \implies Bias increases and Variance decreases

- (a) Increasing the value of λ in lasso regression
 \implies Model complexity decreases
 \implies Bias increases and Variance decreases
- (b) Increasing model complexity by adding more features of high degree
 \implies Bias decreases and Variance increases
- (c) Reducing dimension by choosing only those subset of features which are of more importance
 \implies Model complexity decreases
 \implies Bias increases and Variance decreases

2.

Let the test point be $(x, y = f(x) + \epsilon)$, where $\epsilon \sim \mathcal{N}(0, \sigma^2)$

$$\begin{aligned}
 \text{Mean Squared Error} &= E[(\hat{f}(x) - y)^2] \\
 &= E[\hat{f}(x)^2 + y^2 - 2\hat{f}(x)y] \\
 &= E[(\hat{f}(x) - \overline{\hat{f}(x)})^2] + \overline{\hat{f}(x)}^2 + E[(y - f(x))^2] + f(x)^2 - 2\overline{\hat{f}(x)}f(x) \\
 &\quad \left[\text{Using } E[X^2] = E[(X - \overline{X})^2] + \overline{X}^2; E[\hat{f}(x)] = \overline{\hat{f}(x)}; E[y] = f(x) \right] \\
 &= E[(\hat{f}(x) - \overline{\hat{f}(x)})^2] + (\overline{\hat{f}(x)} - f(x))^2 + \sigma^2 \\
 &= \text{Variance}(\hat{f}(x)) + \text{Bias}(\hat{f}(x))^2 + \text{irreducible noise variance}
 \end{aligned}$$

We can ignore the irreducible noise variance as it can not be reduced even with a perfect estimator. Therefore, Mean Squared Error can be written as sum of $\text{Variance}(\hat{f}(x))$ and $(\text{Bias}(\hat{f}(x)))^2$. Hence proved.

3.2 CS 337: Lab

1.

Code for the function `gen_bais_variance()` updated in notebook.

2.

Code for the function `driver()` updated in notebook.

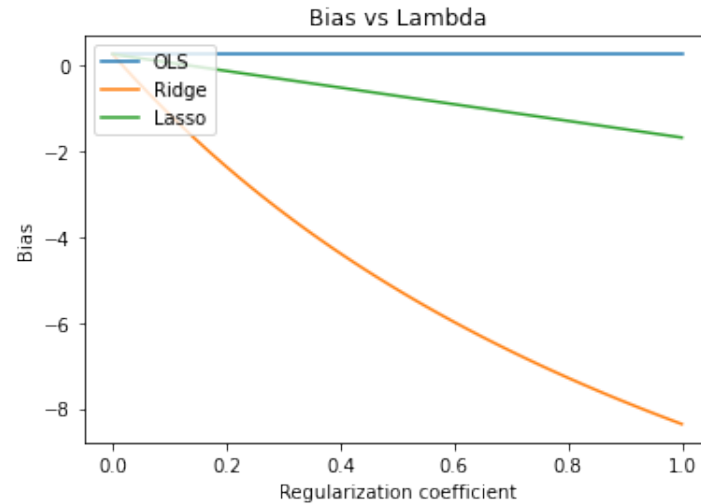


Figure 6: Bias Vs Lambda

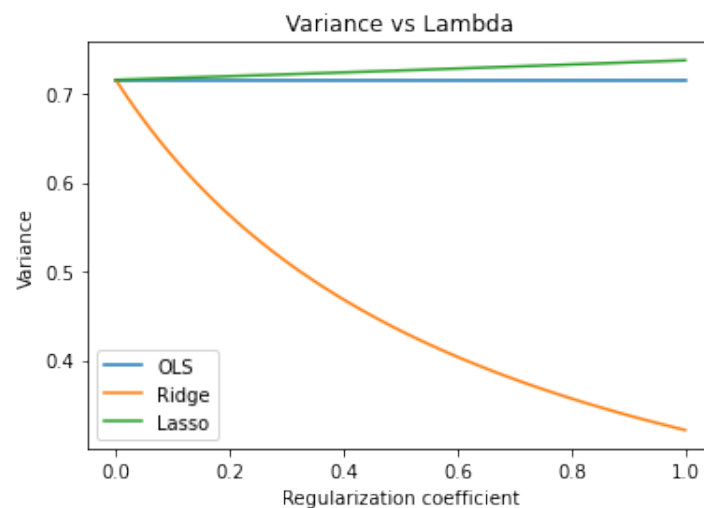


Figure 7: Variance Vs Lambda

Bias is almost zero for OLS, slightly negative for LASSO and significantly negative for Ridge with increase in lambda.

Variance is insignificant w.r.t square of Bias.

On close observation, Variance decreases significantly for Ridge whereas is almost constant for OLS and LASSO.

Overall, OLS has the best performance and it should be as we have a linear data model and linear predicted model.

LASSO underfits slightly as lambda increases because of l_1 constraint on weight.

Ridge underfits significantly as lambda increases because of l_2 constraint on weight.

3.

Code for the functions `ridge()`, `lasso()`, `ols()` updated in notebook.